

AccuKnox AI Security Stack Ranking (Sources from Pg. 2)

Feature / Capability	ACCUKNOX AI Security	WIZ	ORCA SECURITY	PRISMA PROTECT AI	DAXA	HIDDENLAYER	LAKERA	Calypso AI	MINDGARD	Giskard	Prompt:	NOMA	Straiker
AI-SPM / Platform Security	1												
Pipeline Security	2												
Model Security - ML/LLM Scans	3												
Dataset Security	4												
Safety Guardrails - Session Abuse	5												
Safety Guardrails - Unsafe Content	5												
Attack Detection (AI-DR)	6												
Incident Response	7												
Prompt Firewall (inline)	8												
Shadow AI Detection	1												
AIBOM / xBOM	9												
AI Red Teaming	10												
RAG / Dataset Poisoning Protection	11												
Agentic AI Security (Managed Cloud Agents)	12												
MCP Security	12												
AI Gateway Integrations	13												
SDK & Platform Integrations	14												
On Premise Local Models Security	15												
Container Security	16												
Runtime Security (eBPF)	17												

Supported - Capability is GA today, documented on vendor site.
 Partial / Roadmap - Beta, limited scope, or on the roadmap.
 Not Supported - No public evidence of capability.

#	Capability Row	AccuKnox Documentation
1	AI-SPM / Platform Security & Shadow AI Detection	AI-SPM Overview - help.accuknox.com/how-to/aiml-overview/
2	Pipeline Security	CI/CD Integrations - help.accuknox.com/integrations/cicd-overview/
3	Model Security	LLM/ML Static Scans - help.accuknox.com/how-to/llm-static-scan/ , /how-to/ml-static-scan/
4	Dataset Security	Dataset PII/PHI Scans - help.accuknox.com/use-cases/aiml-usecases/
5	Safety Guardrails - Session Abuse & Unsafe Content	Prompt Firewall - help.accuknox.com/use-cases/prompt-firewall-overview/
6	Attack Detection (AI-DR)	AI Detection & Response - help.accuknox.com/use-cases/aidr/ , /use-cases/azure-aidr/
7	Incident Response	CDR & Auto Remediation - help.accuknox.com/getting-started/aws-cdr/
8	Prompt Firewall (inline)	Prompt Firewall Onboard - help.accuknox.com/use-cases/llm-defense-app-onboard/
9	AIBOM / xBOM	xBOM / AIBOM - accuknox.com/platform/xbom
10	AI Red Teaming	AI Red Teaming - help.accuknox.com/use-cases/red-teaming/
11	RAG / Dataset Poisoning Protection	ModelArmor & RAG Fencing - help.accuknox.com/use-cases/modelarmor/
12	Agentic AI / MCP Security	Agentic AI & MCP Server - help.accuknox.com/integrations/mcp-server/
13	AI Gateway Integrations	AI Gateway (Bedrock, Foundry, LiteLLM, Bifrost) - help.accuknox.com/integrations/ai-overview/
14	SDK & Platform Integrations	SDK, Copilot Studio, Power Apps - help.accuknox.com/integrations/ai-overview/
15	On Premise Local Models Security	SaaS vs On-Prem - help.accuknox.com/how-to/aiml-saas-vs-onprem/
16	Container Security	Container Scan - help.accuknox.com/use-cases/container-scan/
17	Runtime Security (eBPF)	Runtime Security (KubeArmor/eBPF) - help.accuknox.com/getting-started/runtime-sec-arch/

1. AccuKnox capability sources

1. AI-SPM and AI Copilot

Multi-cloud AI asset inventory across AWS, Azure, GCP, and on-prem. AskAI copilot answers posture questions and surfaces drift.

- <https://help.accuknox.com/how-to/aiml-overview/>
- <https://help.accuknox.com/how-to/aiml-aws-onboard/>
- <https://help.accuknox.com/how-to/aiml-azure-onboard/>
- <https://help.accuknox.com/how-to/aiml-gcp-onboard/>
- <https://accuknox.com/platform/ai-copilot>

2. AI/ML Pipeline & CI/CD

Pipeline graph view, secrets and IaC scanning, native runners for Jenkins, GitHub Actions, Azure DevOps, Harness, and AWS CodePipeline.

- <https://help.accuknox.com/integrations/cicd-overview/>
- <https://help.accuknox.com/how-to/aiml-overview/>

3. Model Security: ML and LLM Static Scans

Static checks for Pickle, TensorFlow SavedModel, GGUF, DDUF, safetensors, ONNX. Detects deserialization, malicious payloads, supply-chain risk. ModelArmor runtime guards execution.

- <https://help.accuknox.com/how-to/llm-static-scan/>
- <https://help.accuknox.com/how-to/ml-static-scan/>
- <https://help.accuknox.com/use-cases/modelarmor/>

4. Dataset Security (PII / PHI at rest)

PII and PHI scans of training datasets, tenant-custom scan configs, data-poisoning detection on weights and biases.

- <https://help.accuknox.com/use-cases/aiml-usecases/>
- <https://help.accuknox.com/how-to/ml-static-scan/>

5. Safety Guardrails: Session Abuse

Per-session policy enforcement, jailbreak and prompt-injection blocking, PII leak prevention on prompt traffic.

- <https://help.accuknox.com/use-cases/prompt-firewall-overview/>
- <https://help.accuknox.com/use-cases/prompts-categories/>
- <https://accuknox.com/solutions/prompt-firewall>

6. Safety Guardrails: Unsafe Content

Sentiment analysis, hallucination flagging, code-injection blocking on model outputs. Works on cloud-hosted and on-prem LLMs.

- <https://help.accuknox.com/use-cases/prompt-firewall/>
- <https://help.accuknox.com/use-cases/subprompts-categories/>

7. Attack Detection (AI-DR)

CloudTrail and Azure Activity Log ingest, drift detection on AI configs, alerts on public notebooks and shadow models.

- <https://help.accuknox.com/use-cases/aidr/>
- <https://help.accuknox.com/use-cases/azure-aidr/>
- <https://accuknox.com/solutions/ai-dr>

8. Incident Response (CDR, auto-remediation, ITSM)

Cloud Detection and Response on AWS, Azure, GCP. Auto-removes public access. ServiceNow and Jira ticketing.

- <https://help.accuknox.com/getting-started/aws-cdr/>
- <https://help.accuknox.com/getting-started/azure-cdr/>
- <https://help.accuknox.com/getting-started/gcp-cdr/>
- <https://accuknox.com/solutions/ai-dr>

9. Prompt Firewall (inline, gateway, SDK)

Transparent inline proxy, 14 built-in policies plus custom regex, block / sanitize / monitor / allow actions, SDK and gateway modes.

- <https://help.accuknox.com/use-cases/llm-defense-app-onboard/>
- <https://accuknox.com/solutions/prompt-firewall>

10. Shadow AI Detection

Auto-discovery of unsanctioned models, notebooks, and managed AI services across AWS, Azure, GCP, and self-hosted inference engines (Ollama, vLLM).

- <https://help.accuknox.com/how-to/aiml-overview/>
- <https://help.accuknox.com/how-to/aiml-aws-onboard/>

11. AIBOM / xBOM

Generate, sign, and verify SBOM, CBOM, HBOM, AIBOM. Aligned to CERT-In.

- <https://accuknox.com/platform/xbom>

12. AI Red Teaming

1,500+ adversarial probes across four categories. Covers OpenAI, Ollama, custom inference, plus ML model scanning for Pickle, HDF5, TF, ONNX. Findings map to NIST AI RMF, ISO 42001, OWASP LLM Top 10.

- <https://help.accuknox.com/use-cases/red-teaming/>
- <https://accuknox.com/solutions/ai-red-teaming>

13. RAG and Dataset Poisoning Protection

ModelArmor sandboxes model execution, blocks adversarial inputs, and fences PII out of retrieval paths.

- <https://help.accuknox.com/use-cases/modelarmor/>
- <https://help.accuknox.com/use-cases/modelarmor-adversarial-attacks/>
- <https://help.accuknox.com/use-cases/modelarmor-pickle-code/>

14. Agentic AI and MCP Server Security

SPIFFE-based workload identity, OpenFGA fine-grained authorization, MCP server discovery and eBPF-based tool-call sandboxing. ClawArmor isolates OpenClaw instances.

- <https://help.accuknox.com/integrations/mcp-server/>
- <https://accuknox.com/solutions/agentic-ai-security>

15. AI Gateway Integrations

Native connectors for AWS Bedrock AgentCore, Azure AI Foundry / APIM, AWS API Gateway, Apigee, LiteLLM, Bifrost.

- <https://help.accuknox.com/integrations/ai-overview/>
- <https://help.accuknox.com/integrations/bedrock-agentcore/>
- <https://help.accuknox.com/integrations/litellm/>
- <https://help.accuknox.com/integrations/bifrost-integration/>

16. SDK and Platform Integrations

Python SDK for direct app onboarding, Azure Copilot Studio, Microsoft Power Apps integrations.

- <https://help.accuknox.com/integrations/ai-overview/>
- <https://help.accuknox.com/integrations/copilot-studio/>
- <https://help.accuknox.com/integrations/powerapps-integration/>
- <https://help.accuknox.com/use-cases/llm-defense-app-onboard/>

17. On-Premise Local Models Security (AI Factories)

GPU isolation, CUDA gating, kernel-level eBPF runtime protection, policy-driven egress controls for self-hosted and air-gapped deployments. AWS, Azure, Red Hat, Oracle marketplaces.

- <https://help.accuknox.com/how-to/aiml-saas-vs-onprem/>
- <https://help.accuknox.com/getting-started/on-prem-overview/>
- <https://accuknox.com/solutions/securing-ai-factories>

18. Runtime Security (KubeArmor / eBPF)

Zero-trust enforcement at process, file, network, and capability layers via eBPF. No iptables, no kernel modules. K8s, VMs, bare metal, serverless.

- <https://help.accuknox.com/getting-started/runtime-sec-arch/>
- <https://accuknox.com/solutions/securing-ai-factories>

19. Container & Image Security

Registry and image scanning across AWS ECR, Azure ACR, GCR, JFrog, private registries. Vulnerability, secrets, IaC checks.

- <https://help.accuknox.com/use-cases/container-scan/>
- <https://help.accuknox.com/integrations/aws-container-scan/>
- <https://help.accuknox.com/integrations/azure-container-scan/>

2. Competitor verdict justifications

We only list capabilities where the verdict was Not Supported or Partial. A Partial means the vendor offers something adjacent but narrower than the row's intent, or there is no public proof at the depth a specialist would meet. A Not Supported means we found no public evidence on the vendor's own site or product docs.

Wiz

CNAPP leader with strong AI-SPM and AIBOM, weak on inline runtime guardrails and air-gapped deployment.

Capability	Verdict	Evidence and source page
Prompt Firewall	Not Supported	Wiz frames guardrails as an application-layer concern and does not ship a Wiz-branded inline prompt firewall. Their academy article positions guardrails as something customers configure on the application side. https://www.wiz.io/academy/ai-security/llm-guardrails
Safety Guardrails - Unsafe Content	Not Supported	No documented toxic-output or hallucination filtering. AI-SPM verifies Bedrock Guardrails exist; it does not enforce its own. https://www.wiz.io/solutions/ai-spm
AI Gateway Integrations	Not Supported	No named connectors for Bedrock AgentCore, Azure Foundry, LiteLLM, or Bifrost in the integrations directory. https://www.wiz.io/integrations
SDK and Platform Integrations	Not Supported	No Python SDK, Copilot Studio, or Power Apps integration listed. https://www.wiz.io/integrations
On Premise Local Models Security	Not Supported	Wiz is SaaS-only. BYOC is the closest pattern, not air-gapped or self-managed. https://accuknox.com/blog/wiz-ai-spm-alternatives
Safety Guardrails - Session Abuse	Partial	Detects prompt-injection exposure as a posture risk. No inline user-session enforcement. https://www.wiz.io/blog/wiz-ai-spm-secures-ai-agents
Model Security - ML/LLM Scans	Partial	Scans PyTorch and TensorFlow model files. GGUF and safetensors not publicly listed. https://www.wiz.io/blog/wiz-releases-model-scanning
AI Red Teaming	Partial	Red Agent simulates attacks on web apps and APIs. Not a dedicated LLM adversarial-probe suite. https://www.wiz.io/blog/introducing-the-wiz-red-agent
Agentic AI / MCP Security	Partial	Discovers MCP connections and agent identity. No MCP server sandboxing or runtime isolation product. https://www.wiz.io/blog/wiz-ai-spm-secures-ai-agents
Runtime Security (eBPF)	Partial	Wiz Sensor uses eBPF for visibility. Enforcement is not core to the positioning the way it is for

		KubeArmor-class tools. https://www.wiz.io/product/runtime-sensor
--	--	---

Orca Security

Agentless CNAPP with AI-SPM and AIBOM. No inline prompt firewall, no model file static scanning, no air-gapped option.

Capability	Verdict	Evidence and source page
Prompt Firewall	Not Supported	Orca Sensor observes runtime AI activity. There is no inline prompt firewall product. https://orca.security/platform/ai-security/
Safety Guardrails - Unsafe Content	Not Supported	No documented toxic-output or hallucination filtering. https://orca.security/platform/ai-security/
Model Security - ML/LLM Scans	Not Supported	No public evidence of static scanning for Pickle, GGUF, or safetensors model files. https://orca.security/platform/ai-security/
AI Red Teaming	Not Supported	No adversarial-probing or attack-simulation product. https://orca.security/platform/ai-security/
AI Gateway Integrations	Not Supported	No native gateway connectors for Bedrock AgentCore, Azure Foundry, LiteLLM, or Bifrost. https://orca.security/integrations/
SDK and Platform Integrations	Not Supported	No Python SDK, Copilot Studio, or Power Apps integration listed. https://orca.security/integrations/
On Premise Local Models Security	Not Supported	SaaS-only platform. No on-prem or air-gapped option. https://accuknox.com/comparisons/accuknox-vs-orca-security
Pipeline Security	Partial	CI/CD secrets and IaC scanning exist, but they are not framed as AI pipeline security. https://orca.security/resources/blog/infrastructure-as-code-security/
Model Security - ML/LLM Scans (duplicate row label)	Partial	Model inventory and posture exist. There is no runtime model integrity or supply-chain enforcement. https://orca.security/platform/ai-security/
Runtime Security (eBPF)	Partial	Sensor model gives runtime AI activity detection. Sensor is not the eBPF-enforcement category Orca claims. https://orca.security/resources/blog/runtime-ai-security/

Prisma Cloud / Palo Alto Networks AI Stack

Broadest coverage of the CNAPP vendors thanks to AI Runtime Security and the Protect AI acquisition. Weak spot is eBPF positioning and some named gateway integrations.

Capability	Verdict	Evidence and source page
Runtime Security (eBPF)	Partial	Prisma uses an agent plus agentless model. eBPF is not the headline positioning the way it is for KubeArmor-class tools. https://www.paloaltonetworks.com/cloud-security/runtime-security
Attack Detection (AI-DR)	Partial	AIRS 3.0 adds shadow AI discovery. AI-DR is less explicit as a category than at the specialists. https://www.paloaltonetworks.com/blog/2026/03/prisma-air-3-0-autonomous-ai/
AI Gateway Integrations	Partial	Confirmed integrations for LiteLLM guardrail and Copilot Studio. Azure Foundry on the roadmap. No Bedrock AgentCore or Bifrost connector documented. https://docs.litellm.ai/docs/proxy/guardrails/panw_prisma_air
AIBOM / xBOM	Partial	Protect AI Guardian generates AIBOM. It is not heavily promoted inside unified Prisma AIRS marketing yet. https://protectai.com/guardian
Dataset Security	Supported	DSPM has 100+ classifiers across S3, Azure Data Lake, Snowflake. https://www.paloaltonetworks.com/cyberpedia/what-is-dspm

DAXA.ai

Pebblo-based safe-prompt / safe-inference / RAG protection. Strong on data layer, no static model scanning, no container or runtime workload security.

Capability	Verdict	Evidence and source page
Pipeline Security	Not Supported	No CI/CD secrets or IaC scanning documented. https://daxa.ai/pebblo
Container Security	Not Supported	DAXA does not ship a container image scanner. https://daxa.ai/product
Runtime Security (eBPF)	Not Supported	No eBPF runtime enforcement layer. https://daxa.ai/product
Model Security - ML/LLM Scans	Not Supported	No static scanning of Pickle, GGUF, or safetensors files. https://daxa.ai/pebblo
AI Red Teaming	Not Supported	No adversarial red-team probe suite documented. https://daxa.ai/product
Automated Cloud Remediation	Not Supported	Slack and Jira workflows only. No cloud auto-remediation. https://daxa.ai/product
Attack Detection (AI-DR)	Partial	Auditability and unified visibility over AI activity. Not cloud-log-based AI-DR. https://daxa.ai/product
Compliance Mapping	Partial	EU AI Act, HIPAA, GDPR, SOC 2, PCI mentioned. MITRE ATLAS and NIST AI RMF not explicitly mapped. https://daxa.ai/product

Protect AI (now part of Palo Alto)

Strong on model artefact scanning (Guardian, ModelScan), runtime app guardrails (Layer), and red teaming (Recon). No dataset PII at rest. AIBOM is light.

Capability	Verdict	Evidence and source page
Dataset Security	Not Supported	Focus is on models, not dataset PII at rest. https://protectai.com/guardian
Automated Cloud Remediation	Not Supported	Integrations with PagerDuty, DataDog, Splunk, Elastic. No native ITSM auto-remediation. https://protectai.com/layer
Container Security	Partial	Guardian ships as a Docker container that scans models. Not arbitrary container image scanning. https://protectai.com/guardian
Shadow AI Detection	Partial	Layer auto-discovers AI apps. Cross-cloud shadow model discovery is limited. https://protectai.com/layer
AIBOM / xBOM	Partial	Model metadata, sources, findings tracked. Not a formal AIBOM or xBOM standard output. https://protectai.com/guardian
Agentic AI / MCP Security	Partial	Layer tracks tool calls, function calls, multi-turn agent flows. No documented MCP-specific controls. https://protectai.com/layer
AI Gateway Integrations	Partial	AWS Bedrock integration confirmed. Azure Foundry, LiteLLM, Bifrost not documented. https://protectai.com/layer
On Premise Local Models Security	Partial	Docker container model and modular deployment. Full air-gapped support not stated. https://protectai.com/guardian

HiddenLayer

Specialist on model artefact security plus AI runtime security. Strong on model scanning, AIBOM, and red teaming. Weak on dataset, container, and eBPF runtime.

Capability	Verdict	Evidence and source page
Dataset Security	Not Supported	No dedicated dataset PII or PHI scanner. https://hiddenlayer.com/aisec-platform/
Container Security	Not Supported	HiddenLayer is not a container image scanner. https://hiddenlayer.com/aisec-platform/
Runtime Security (eBPF)	Not Supported	No eBPF kernel enforcement. Runtime sits outside the inference path. https://www.hiddenlayer.com/platform/ai-runtime-security
Incident Response	Partial	SIEM and SOAR integration. No native ITSM workflow. https://docs.hiddenlayer.ai/docs/products/console/runtime_protection_interactions
Safety Guardrails - Unsafe Content	Partial	Guardrails exist. Less emphasis on

		toxicity or output safety vs injection focus. https://docs.hiddenlayer.ai/docs/products/console/runtime_protection_interactions
AI Gateway Integrations	Partial	Generic API gateway support. No named Bedrock, Foundry, or LiteLLM connectors documented. https://hiddenlayer.com/platform/ai-runtime-security/
On Premise Local Models Security	Partial	Agentless lightweight containers. On-prem mentioned. Air-gap not the headline story. https://hiddenlayer.com/aisec-platform/

Lakera (now Cisco AI Defense)

Specialist guardrail and red-team vendor. Strong on prompt firewall, agent MCP, and red teaming. Not a CNAPP, so no asset inventory, no container security, no AIBOM.

Capability	Verdict	Evidence and source page
AI-SPM / Platform Security	Not Supported	Lakera is not a cloud posture or AI inventory product. https://www.lakera.ai/
Pipeline Security	Not Supported	No CI/CD secrets or IaC scanning. https://www.lakera.ai/
Model Security - ML/LLM Scans	Not Supported	Runtime guardrails only. No model supply-chain checks. https://www.lakera.ai/
Dataset Security	Not Supported	No dataset-at-rest PII scanner. https://www.lakera.ai/
Container Security	Not Supported	Not a container scanner. https://www.lakera.ai/
Runtime Security (eBPF)	Not Supported	API-based. No eBPF runtime. https://www.lakera.ai/
AIBOM / xBOM	Not Supported	No AIBOM feature. https://www.lakera.ai/
Incident Response	Not Supported	Detection only. No remediation automation. https://www.lakera.ai/
Attack Detection (AI-DR)	Partial	Workforce AI Security covers shadow AI. No cloud-log-based AI-DR. https://www.lakera.ai/products/workforce-ai-security
AI Gateway Integrations	Partial	Model-agnostic. Works in front of Bedrock, Azure OpenAI. No named LiteLLM, Foundry, or Bifrost connectors. https://docs.lakera.ai/

CalypsoAI (now F5 AI Guardrails / AI Red Team)

Guardrails and red-teaming, sold as part of the F5 ADSP. No CNAPP coverage, no model file scanning, no AIBOM, no shadow AI.

Capability	Verdict	Evidence and source page
AI-SPM / Platform Security	Not Supported	No multi-cloud AI inventory or AI-SPM. https://www.f5.com/products/ai-guardrails

Pipeline Security	Not Supported	No CI/CD secrets or IaC scanning. https://www.f5.com/products/ai-guardrails
Dataset Security	Not Supported	Runtime DLP only. Not dataset-at-rest scanning. https://www.f5.com/products/ai-guardrails
Container Security	Not Supported	No container image scanning. https://www.f5.com/products/ai-guardrails
Runtime Security (eBPF)	Not Supported	No eBPF runtime enforcement. https://www.f5.com/products/ai-guardrails
Shadow AI Detection	Not Supported	No shadow AI inventory feature documented. https://www.f5.com/products/ai-guardrails
Model Security - ML/LLM Scans	Not Supported	No Pickle, GGUF, or safetensors static scanning. https://www.f5.com/products/ai-guardrails
AIBOM / xBOM	Not Supported	No AIBOM. https://www.f5.com/products/ai-guardrails
RAG / Dataset Poisoning	Not Supported	No RAG or dataset poisoning protection. https://www.f5.com/products/ai-guardrails
Incident Response	Partial	Guardrails feeds Red Team findings. No native ITSM described. https://www.f5.com/products/ai-red-team

Mindgard

AI red teaming flagship. Strong on automated probes and shadow AI. Light on CNAPP-style coverage and AIBOM.

Capability	Verdict	Evidence and source page
Dataset Security	Not Supported	No dataset-at-rest PII or PHI scanning. https://mindgard.ai/ai-security-platform
Container Security	Not Supported	No container image scanning. https://mindgard.ai/ai-security-platform
Runtime Security (eBPF)	Not Supported	No eBPF enforcement. https://mindgard.ai/ai-security-platform
AIBOM / xBOM	Not Supported	No AIBOM. https://mindgard.ai/ai-security-platform
On Premise Local Models Security	Not Supported	SaaS-only documented. No air-gapped deployment. https://mindgard.ai/ai-security-platform
Attack Detection (AI-DR)	Partial	Shadow AI discovery via recon. No cloud-log AI-DR. https://mindgard.ai/ai-security-platform
Prompt Firewall	Partial	Runtime detection and response with context-driven guardrails. Less mature than dedicated firewalls. https://mindgard.ai/ai-security-platform
Incident Response	Partial	Remediation guidance. No ITSM ticketing documented.

		https://mindgard.ai/ai-security-platform
--	--	---

Giskard

Pre-deployment LLM testing and red-teaming framework. Strong on red teaming, compliance mapping, and SDK. Not built for runtime, infra, or cloud posture.

Capability	Verdict	Evidence and source page
AI-SPM / Platform Security	Not Supported	No cloud AI inventory or AI-SPM. https://www.giskard.ai/
Pipeline Security	Not Supported	No CI/CD secrets or IaC scanning. https://www.giskard.ai/
Dataset Security	Not Supported	No dataset-at-rest scanner. https://www.giskard.ai/
Container Security	Not Supported	No container scanning. https://www.giskard.ai/
Runtime Security (eBPF)	Not Supported	Testing platform, not eBPF runtime. https://www.giskard.ai/
Prompt Firewall	Not Supported	Testing platform, not inline runtime firewall. https://www.giskard.ai/
Shadow AI Detection	Not Supported	No shadow AI inventory. https://www.giskard.ai/
Model Security - ML/LLM Scans	Not Supported	Model scanner covers ML behaviour. It does not scan Pickle, GGUF, or safetensors files. https://www.giskard.ai/products/llm-evaluation-hub
AIBOM / xBOM	Not Supported	No AIBOM. https://www.giskard.ai/
Attack Detection (AI-DR)	Not Supported	Not an AI-DR product. https://www.giskard.ai/
Incident Response	Not Supported	No incident response workflow. https://www.giskard.ai/
AI Gateway Integrations	Not Supported	No gateway connectors. https://www.giskard.ai/

Prompt Security

Runtime prompt firewall plus MCP gateway with shadow AI discovery. Strong on guardrails, weak on model artefact, dataset, container, and AIBOM.

Capability	Verdict	Evidence and source page
AI-SPM / Platform Security	Not Supported	Focus is prompts and agents, not cloud AI asset inventory. https://www.prompt.security/
Model Security - ML/LLM Scans	Not Supported	No model file static scanning. https://www.prompt.security/
Container Security	Not Supported	Not a container scanner. https://www.prompt.security/
Runtime Security (eBPF)	Not Supported	Proxy / gateway. Not eBPF. https://www.prompt.security/
AIBOM / xBOM	Not Supported	No AIBOM. https://www.prompt.security/
Pipeline Security	Partial	Code assistant guardrails catch secrets in the IDE. No IaC scanning. https://www.prompt.security/
Dataset Security	Partial	DLP on prompts and responses. No

		dataset-at-rest scanning. https://www.prompt.security/
Incident Response	Partial	SIEM and ticketing integrations. No cloud auto-remediation. https://www.prompt.security/
AI Gateway Integrations	Partial	LLM-agnostic proxy fronts Bedrock, Azure, OpenAI. No native LiteLLM or Bifrost plugins documented. https://www.prompt.security/

Noma Security

MLOps and AI agent security platform. Closest to AccuKnox in breadth. Gaps: full eBPF runtime, named cloud auto-remediation playbooks, and named gateway connectors beyond LiteLLM and Copilot Studio.

Capability	Verdict	Evidence and source page
Pipeline Security	Partial	CI/CD and SCM integrations. Limited public detail on secrets and IaC scanning. https://noma.security/platform/
Container Security	Partial	Containers referenced in NanoClaw context. No dedicated image scanner GA. https://noma.security/platform/
Runtime Security (eBPF)	Partial	Runtime guardrails on prompts and tools. No public eBPF sensor advertised. https://noma.security/platform/runtime-protection/
Attack Detection (AI-DR)	Partial	Runtime observability across prompts and agents. No explicit cloud-log AI-DR module. https://noma.security/platform/aispm/
Incident Response	Partial	Audit trails, policy enforcement, ServiceNow integration. Limited auto-remediation detail. https://noma.security/platform/runtime-protection/
RAG / Dataset Poisoning	Partial	Data-poisoning checks for agents. RAG poisoning not separately branded. https://noma.security/platform/
AI Gateway Integrations	Partial	Copilot Studio, Salesforce Agentforce, LiteLLM guardrail plugin documented. Bedrock and Bifrost not confirmed as native. https://docs.litellm.ai/docs/proxy/guardrails/noma_security

Straiker.ai

Agentic-AI security specialist. Strong on Ascend (red teaming), Defend (runtime guardrails), and MCP. SaaS-only, no model file scanning, no AIBOM, no air-gapped deployment.

Capability	Verdict	Evidence and source page
Model Security - ML/LLM Scans	Not Supported	Focus is agent runtime, not model artefact protection or supply chain. https://www.straiker.ai/products/discover-ai

Container Security	Not Supported	No container image scanning offering. https://www.straiker.ai/
LLM/ML Static Scans	Not Supported	No Pickle, GGUF, or safetensors static scanning advertised. https://www.straiker.ai/
AIBOM / xBOM	Not Supported	No AIBOM or xBOM feature. https://www.straiker.ai/
On Premise Local Models Security	Not Supported	SaaS-only. Inline gateway and sensors. No air-gapped offering documented. https://www.straiker.ai/products/defend-ai
AI-SPM / Platform Security	Partial	Discover AI gives Agent-SPM and AI-SPM. Agent-centric, not full multi-cloud platform inventory. https://www.straiker.ai/products/discover-ai
Pipeline Security	Partial	Ascend tests in CI/CD pipelines. No dedicated secrets or IaC scanner. https://www.straiker.ai/products/ascend-ai
Dataset Security	Partial	Defend prevents PII, PCI, and HIPAA exfiltration at runtime. No at-rest dataset scanning. https://www.straiker.ai/products/defend-ai
Runtime Security (eBPF)	Partial	eBPF sensors mentioned for runtime telemetry. Not a general workload eBPF enforcer. https://www.straiker.ai/products/defend-ai
Shadow AI Detection	Partial	Discover AI catalogs agents and MCP servers. Agent-shadow focus rather than unsanctioned models. https://www.straiker.ai/products/discover-ai
Incident Response	Partial	Alerts and audit logs. Limited ITSM or auto-remediation detail. https://www.straiker.ai/products/defend-ai